

HW2 STA521 Fall18

Eduardo Coronado - ec243 - ecoronado92

Due September 23, 2018 5pm

Background Reading

Readings: Chapters 3-4 in Weisberg Applied Linear Regression

Exploratory Data Analysis

0. Preliminary read in the data. After testing, modify the code chunk so that output, messages and warnings are suppressed.

1. **Create a summary of the data. How many variables have missing data? Which are quantitative and which are qualitative?**

From the summary data below 6 out of the 7 variables have missing data, with **ModernC** and **Frate** being the ones with most NAs present. Also from the metadata we know that only 125 observations out of 210 have complete data for all variables.

##	ModernC	Change	PPgdp	Frate
##	Min. : 1.00	Min. : -1.100	Min. : 90	Min. : 2.00
##	1st Qu.: 19.00	1st Qu.: 0.580	1st Qu.: 479	1st Qu.: 39.50
##	Median : 40.50	Median : 1.400	Median : 2046	Median : 49.00
##	Mean : 38.72	Mean : 1.418	Mean : 6527	Mean : 48.31
##	3rd Qu.: 55.00	3rd Qu.: 2.270	3rd Qu.: 8461	3rd Qu.: 58.00
##	Max. : 83.00	Max. : 4.170	Max. : 44579	Max. : 91.00
##	NA's : 58	NA's : 1	NA's : 9	NA's : 43

##	Pop	Fertility	Purban
##	Min. : 2.3	Min. : 1.000	Min. : 6.00
##	1st Qu.: 767.2	1st Qu.: 1.897	1st Qu.: 36.25
##	Median : 5469.5	Median : 2.700	Median : 57.00
##	Mean : 30281.9	Mean : 3.214	Mean : 56.20
##	3rd Qu.: 18913.5	3rd Qu.: 4.395	3rd Qu.: 75.00
##	Max. : 1304196.0	Max. : 8.000	Max. : 100.00
##	NA's : 2	NA's : 10	

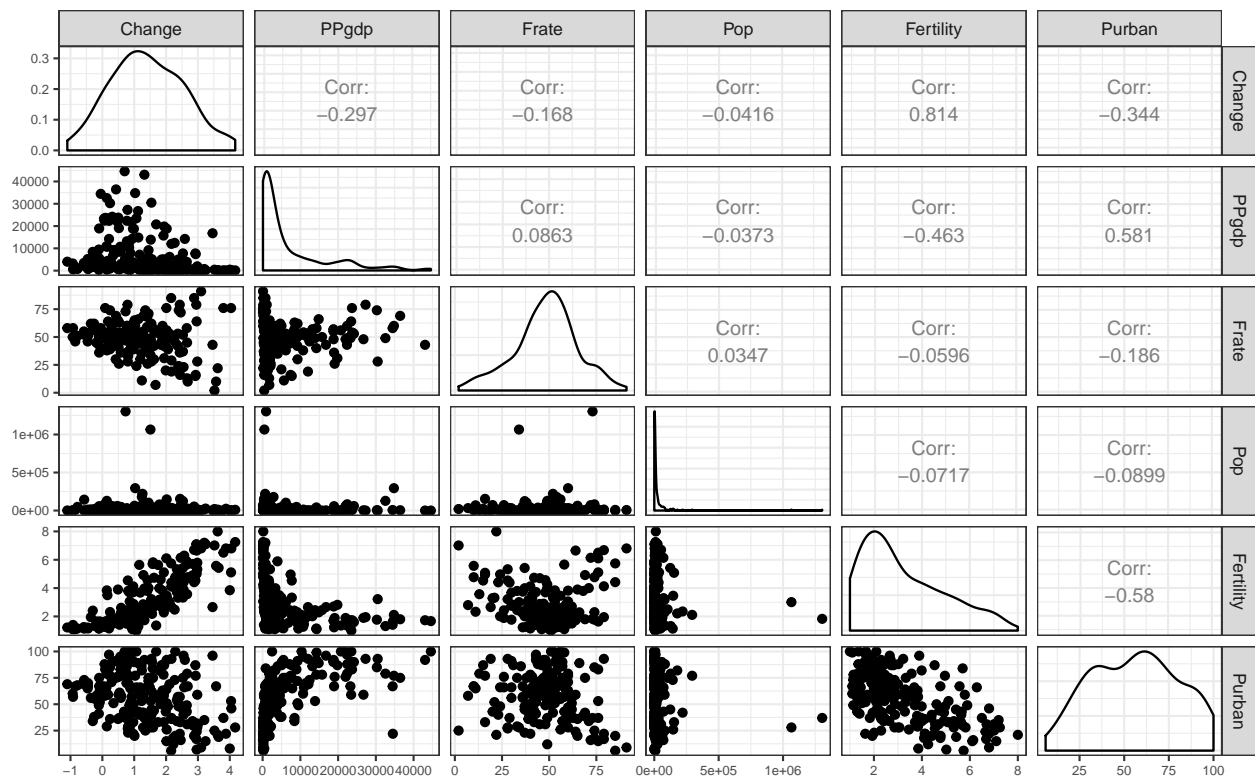
Additionally, **all variables are quantitative.**

```
## 'data.frame': 210 obs. of 7 variables:
## $ ModernC : int NA NA 49 NA NA NA 51 NA 22 NA ...
## $ Change : num 3.88 0.68 1.67 2.37 2.59 3.2 0.53 1.17 -0.45 2.02 ...
## $ PPgdp : int 98 1317 1784 NA 14234 739 8461 7163 687 NA ...
## $ Frate : int NA NA 7 42 NA NA 63 44 51 53 ...
## $ Pop : num 23897 3167 31800 57 64 ...
## $ Fertility: num 6.8 2.28 2.8 NA NA 7.2 NA 2.44 1.15 NA ...
## $ Purban : int 22 43 58 53 92 35 37 88 67 51 ...
```

2. What is the mean and standard deviation of each quantitative predictor? Provide in a nicely formatted table

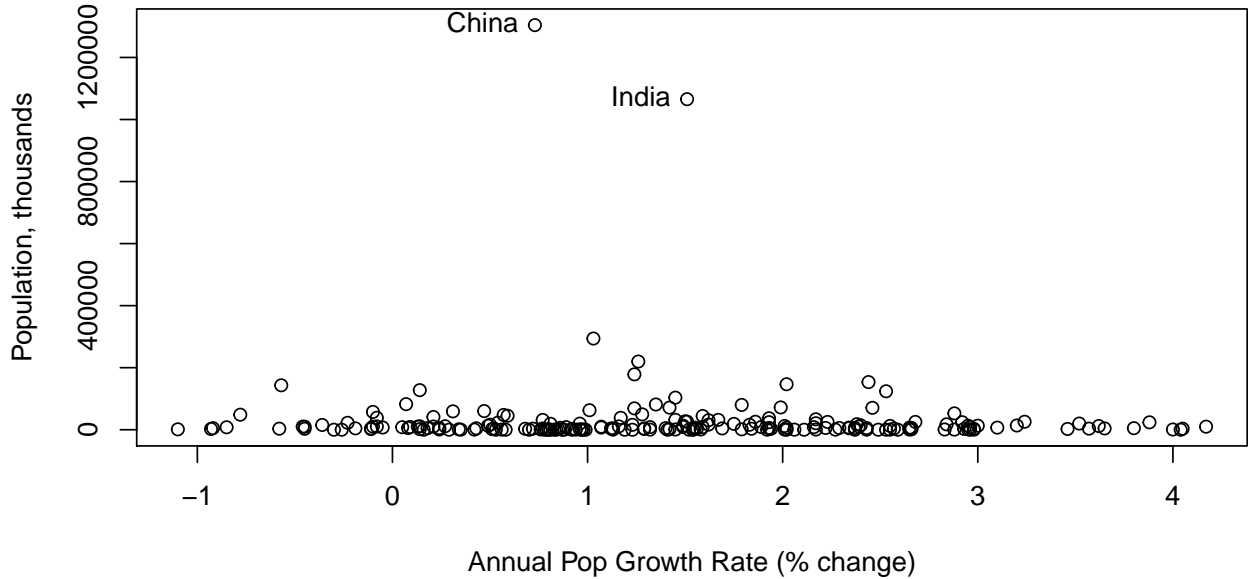
	Mean	Standard Dev
ModernC	38.72	22.64
Change	1.42	1.13
PPgdp	6527.39	9325.19
Frate	48.31	16.53
Pop	30281.87	120676.69
Fertility	3.21	1.71
Purban	56.20	24.11

3. Investigate the predictors graphically, using scatterplots or other tools of your choice. Create some plots highlighting the relationships among the predictors. Comment on your findings regarding trying to predict 'ModernC' from the other variables. Are there potential outliers, nonlinear relationships or transformations that appear to be needed based on your graphical EDA?



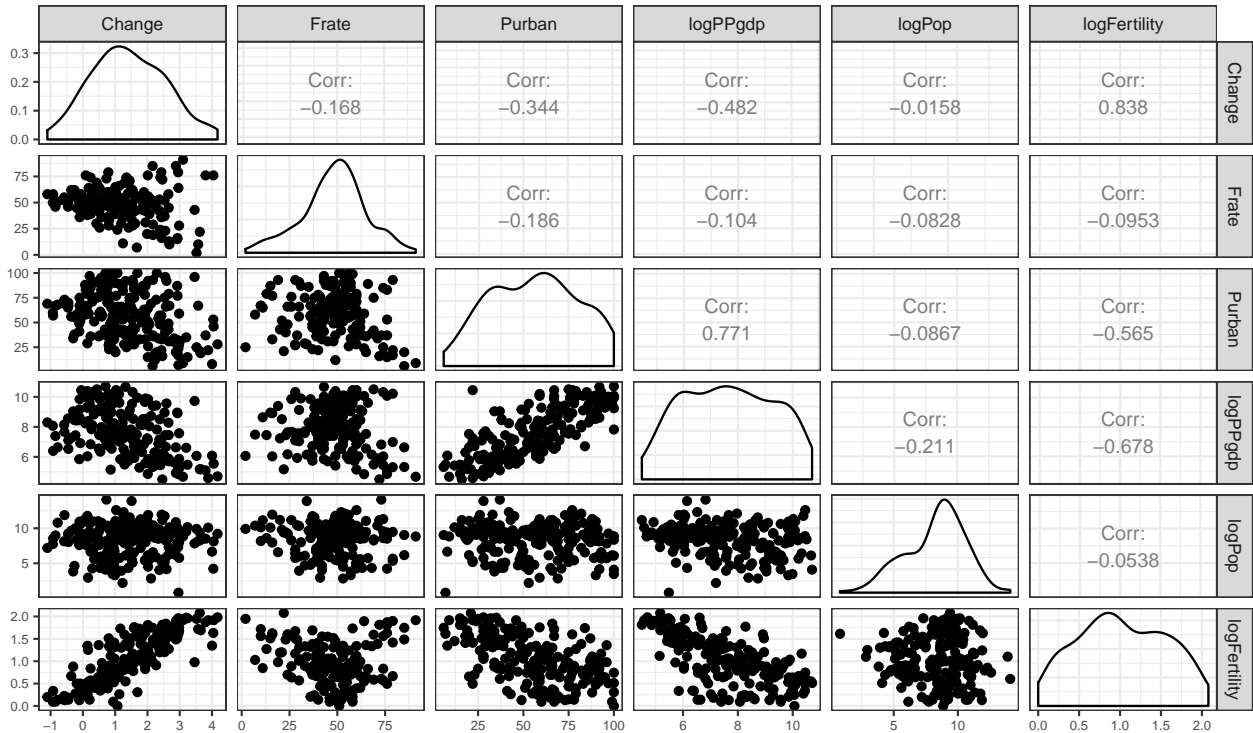
Using the `pairs` function we can do an initial assessment of the relationships among the predictor values. From **Fig 1** above it is easy to notice at first glance that many relationships among predictor variables seem non-linear. The relationships between the **Change**, **Purban** and **Fertility** variables seem to be the ones that mostly resemble a somewhat linear relationship. The **Frate** variable does seem to have a non-multicollinear relationship with some of the variables. However, two plots stand out from this plot - **PPgdp** and **Pop**. The **PPgdp** predictor's relationships seem to follow an increasing or decreasing exponential, while the **Pop** helps signal what seem to be two clear outliers (**Fig2**, below). From this plot we can see that China and India seem as outliers in terms of population vs other countries.

Fig 2. Annual Pop Growth Rate Change vs Population for 210 Countries



Nevertheless, it is important to note that the scale of these two variables is several orders of magnitude higher than the predictor they are compared against, which brings to mind possible linear transformations for further evaluations. Thus, I explored whether a simple $\log()$ transformation would suffice to demonstrate a linear relationship among the non-linear relationships in **Fig 1** (i.e. $\log\text{Fertility}$, $\log\text{PPgdp}$, $\log\text{Pop}$). Even though this was a crude first transformation on the date, in **Fig 3** below we can notice that transforming the data can help elucidate potential linear relationships among our predictors, as well as the presence or lack of multicollinearity.

Fig 3. Pairwise Comparisons of Transformed Predictor Variables from UN3 Dataset



Finally, we can see that a linear combination of the predictor variables could be helpful to predict `ModernC` given these seem to follow a linear relationship, some with would need transformations and others without transformations. However, it is important to note that some predictors exhibit multicollinearity which means adding them to the linear model would be redundant as these would be contributing the same variance to the response variable. It is something that we should consider as we continue to build and assess the fit of the model.

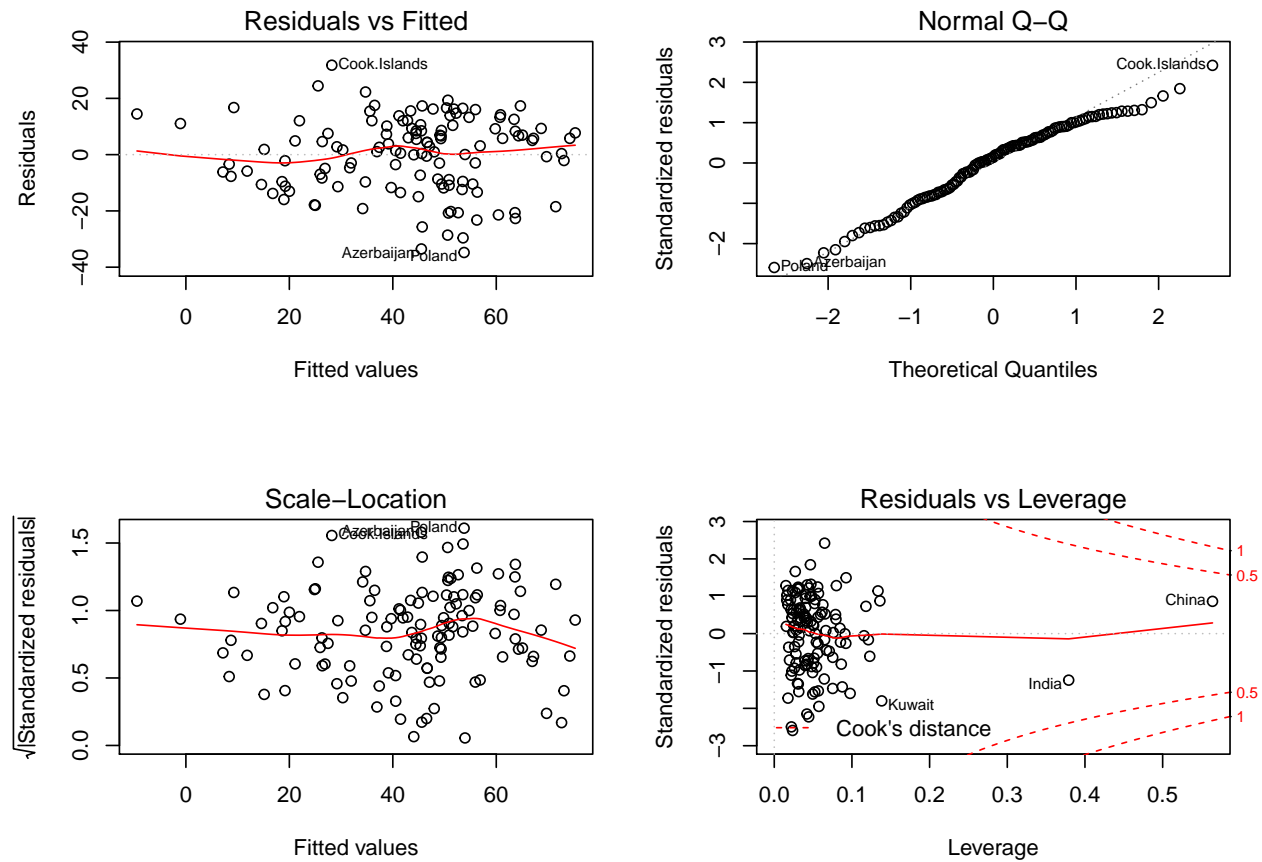
4. Use the `lm()` function to perform a multiple linear regression with `ModernC` as the response and all other variables as the predictors, using the formula `ModernC ~ .`, where the `.` includes all remaining variables in the dataframe. Create diagnostic residual plot from the linear model object and comment on results regarding assumptions. How many observations are used in your model fitting?

From the initial comparisons in **Fig3** we know that the predictor variables have a linear relationship which allows us to interpret these multiple regression diagnostic plots as those of a simple regression model.

***Note:** the above linear model and diagnostic plots were done using the original, non-transformed data*

Using the `summary` function, we notice that the `lm` function automatically excluded 85 observations. Therefore, the multiple regression model was done 125 observations. Looking at **Fig 5** we can notice a minor heteroscedastic trend on the top-left plot, which shows that the variances is non-constant. However, this trend isn't that significant thus our assumptions of linearity still hold. On the Normal Q-Q plot we observe some points diverging from the normal line - especially on the top-right - which means our data follows a skewed normal distribution. Yet, our linearity assumption still holds as the observed standard deviations seem to follow the theoretical ones. Similarly, show that the errors for some observations are above 1 standard deviation and there are observations having a higher influence on the fit (i.e. China and India). However, these aren't significant enough to discard our assumptions of linearity. Overall, Cook Islands, Azerbaijan, Poland, China, and India are likely candidates for outlier testing given these observations tend to be farther away from the rest of the data or are highly influential.

Fig 5. Residual Plots for a Linear Model Fit of ‘ModernC’



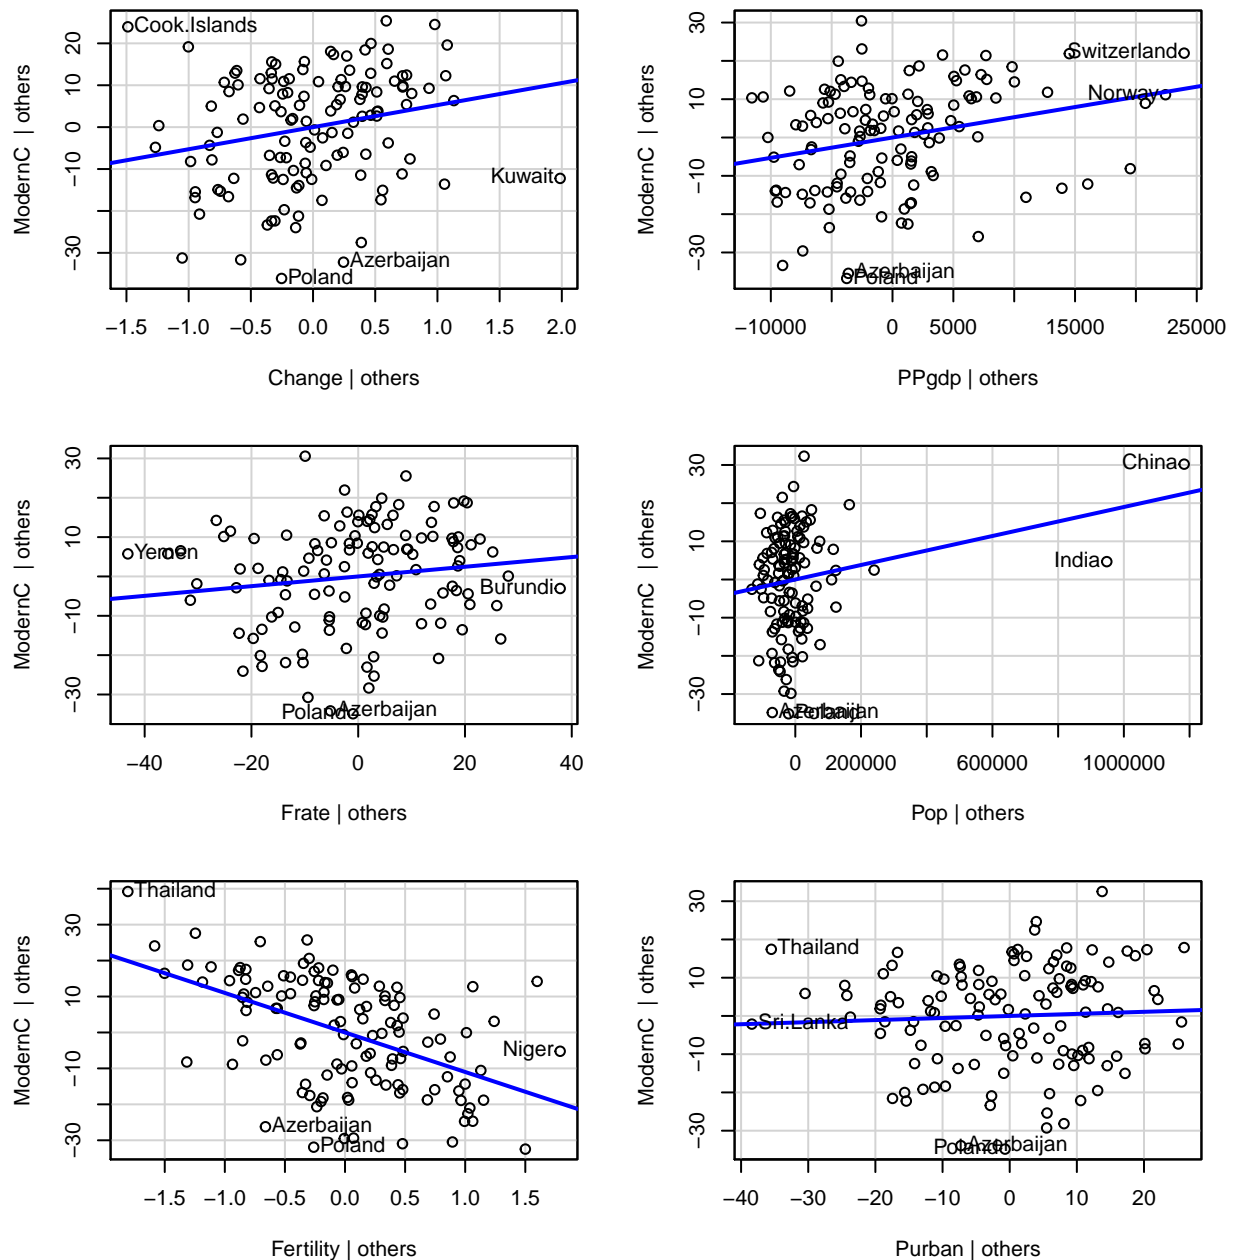
5. Examine added variable plots ‘car::avPlot’ or ‘car::avPlots’ for your model above. Are there any plots that suggest that transformations are needed for any of the terms in the model? Describe. Is it likely that any of the localities are influential for any of the terms? Which localities? Which terms?

From the plots in Fig 6 we can observe that a transformation on the Pop term would be helpful. This is noticeable by the large amount of data points concentrated around zero while only China and India seem to be spread out. Adding a linear transformation, such as log, would help to reduce the skewedness of the plot by expanding those values near zero and contracting those away from zero. It could also help reduce the current influence of the China/India observations on the fit of the model.

It is also noticeable how certain countries are influential for specific terms. For example, Kuwait and Cook’s Island are an influential point for **Change**. India and China, as seen before, are influential on the **Pop** term. Noticeably, these plots brought up new influential localities for specific terms such as Norway and Switzerland for **PPgdp**, Niger and Thailand for **Fertility**, Yemen and Burundi for **Frate**, and Sri Lanka and Thailand for **Purban**. Again, we can notice previous localities such as Poland, Azerbaijan, and Cook’s Island as being influential for certain terms such as **PPgdp** or **Purban**, among others.

Overall, from these plots we can notice the explanatory power of each predictor on the response variable after accounting for all the other predictors. From the slope we can notice almost all terms have either a positive or negative linear relationship with the response variable and thus contribute to explain the variability. **Purban** and **Frate** seem to have the least explanatory power after accounting for all other predictors. A possible cause for this would be an existing multi-colinear relationship with another term already accounted in the linear model.

Fig 6. Added-Variable Plots for 6 Predictors of 'modern_Im' Model




6. Using the Box-Tidwell 'car::boxTidwell' or graphical methods find appropriate transformations of the predictor variables to be used as predictors in the linear model. If any predictors are negative, you may need to transform so that they are non-negative. Describe your method and the resulting transformations.

As mentioned in the previous question, Pop seemed as the most promising candidate that would benefit from a transformation given the data was clustered near zero with two points spread far away. Using the `boxTidwell` built-in function the optimal transformation is close to $\lambda = 0.5$ (i.e. \sqrt{Pop}). However, from this method the transformation does not significantly improve the fit as shown by a non-significant p-value (i.e. the transformation isn't significantly different from $\lambda = 1$ - our H_0). I similarly tested other predictors, but the suggested transformations didn't provide enough evidence to reject the H_0 . Just to double check that highly influential points on Pop such as China and India weren't giving us a different transform estimation, I

removed them from a modified dataframe and re-test. Even removing this we fail to reject H_0 . However, graphically from **Fig 3** we can see that a `log()` transformation does improve the linear relationship between the `Pop`, `PPgdp` and `ModernC` - thus I will use these transformations.

7. Given the selected transformations of the predictors, select a transformation of the response using ‘MASS::boxcox’ or ‘car::boxCox’ and justify.

From the plot below we can observe the MLE for λ using `boxCox`. Since the 95% confidence interval doesn’t include 1 and the `boxCox` $\lambda = 0.76$ I decided to use the closest λ that would still allow for interpretable results $\lambda = 0.5$. This was train of thought is justified given we want a model that improved by a transformation, however we don’t want it to be complicated when interpreting results.



8. Fit the regression using the transformed variables. Provide residual plots and added variables plots and comment. If you feel that you need additional transformations of either the response or predictors, repeat any steps until you feel satisfied.

From the previous question I had decided to use $\lambda = 0.5$ for the response variable given it was close to the significant MLE for this parameter. However, after plotting the residual and added variable plots I noticed this transformation wasn’t beneficial. After trying the λ ’s MLE and $\lambda = 1$, I decided to proceed with the later given it improves the model and interpretability. Thus in **Fig 7** we observe the residual plots for,

$$\text{ModernC} = \beta_0 + \beta_1 \text{Change} + \beta_2 \log \text{PPgdp} + \beta_3 \text{Frate} + \beta_4 \log \text{Pop} + \beta_5 \text{Frate} + \beta_6 \text{Purban}$$

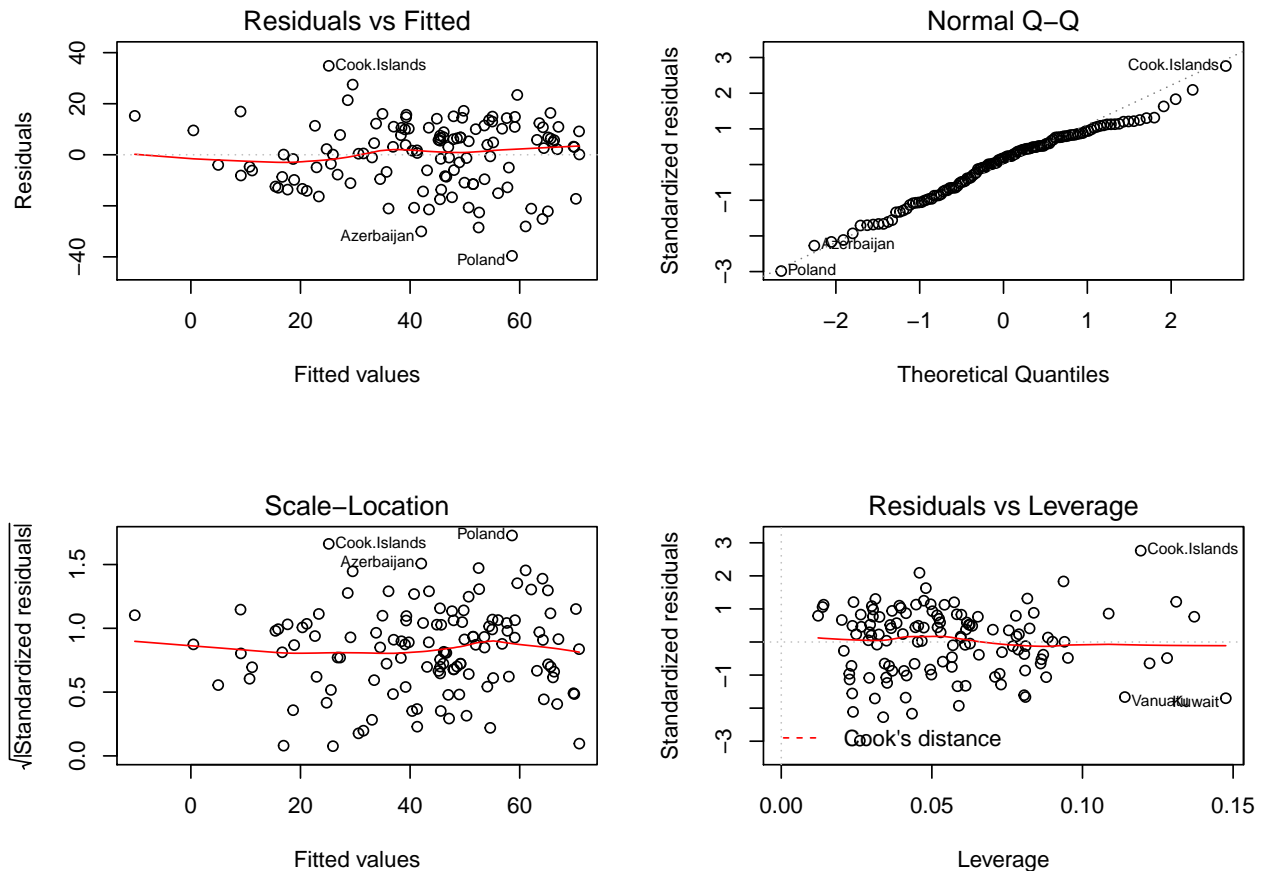
Compared to **Fig 5**, we can observe two improvements: 1) we don’t see the previous high leverage points (China/India) near a Cook’s Distance that would raise a red flag for influential points nor any other point, 2) there is a minor correction of the points in both tails toward the diagonal line in the Normal Q-Q. Even when the transformations did very slight corrections to the minor heteroscedastic trends, these are still within reasonable bounds.

Fig 7. Residual Plots for a Transformed Linear Model Fit of ‘ModernC’

```
##
## Call:
## lm(formula = ModernC ~ Change + logPPgdp + Frate + logPop + Fertility +
##      Purban, data = UN3_log[, c(-3, -5, -10)])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -39.597  -9.540   2.238  10.024  34.840
```

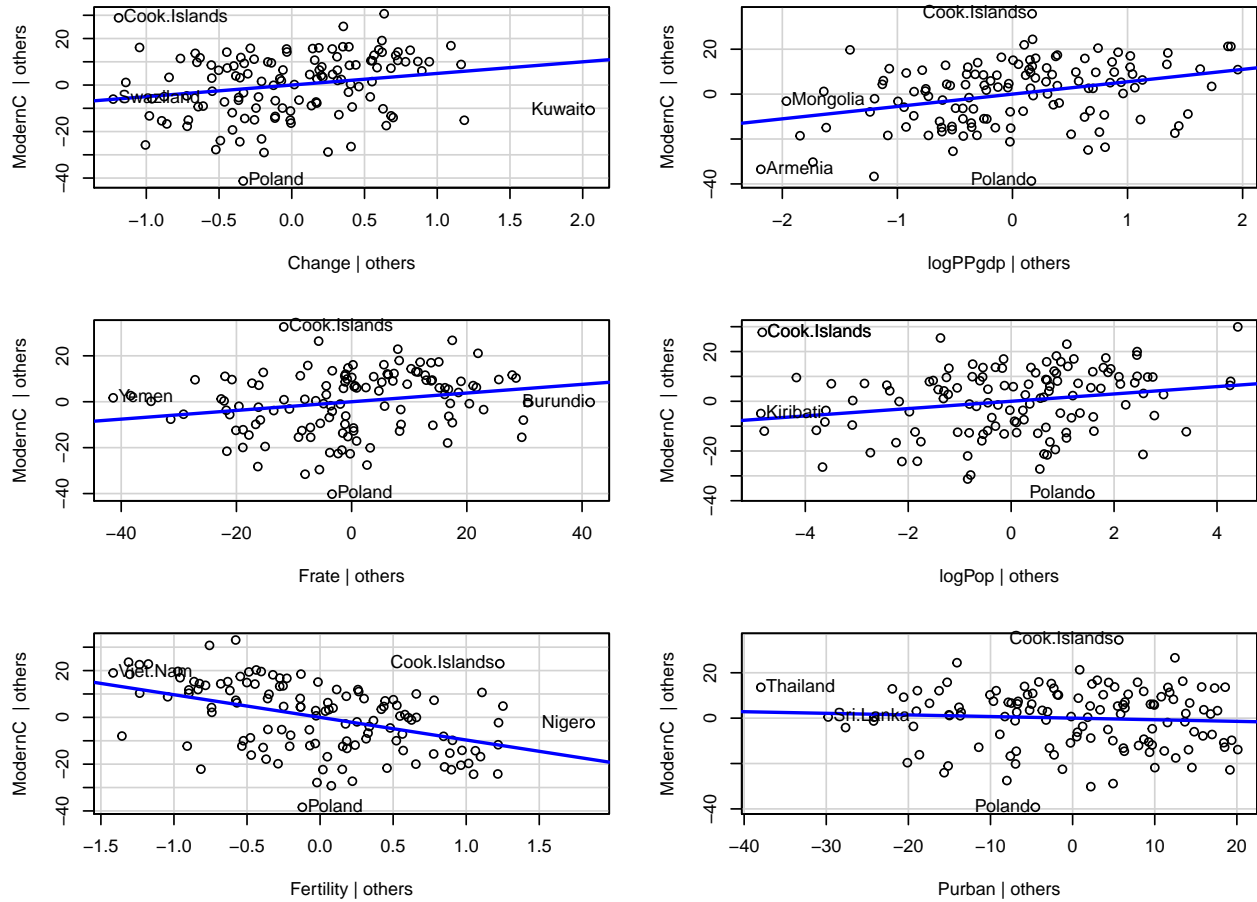


```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.11547    14.50854   0.284 0.777169
## Change       4.99296     2.07709   2.404 0.017781 *
## logPPgdp     5.50728     1.40505   3.920 0.000149 ***
## Frate        0.18939     0.07711   2.456 0.015500 *
## logPop       1.47207     0.62875   2.341 0.020897 *
## Fertility    -9.67594     1.76561  -5.480 2.44e-07 ***
## Purban      -0.07077     0.09760  -0.725 0.469829
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.44 on 118 degrees of freedom
## (85 observations deleted due to missingness)
## Multiple R-squared:  0.626, Adjusted R-squared:  0.6069
## F-statistic: 32.91 on 6 and 118 DF,  p-value: < 2.2e-16
```




Comparing the added-variable plots, we can see that in **Fig 8** two changes on the transformed predictors compared to the plots of **Fig 6**. \logPop plot's slope is less prominent than without a transformation now that we have China or India closer to the other observations and not dictating the value of the coefficient. In \logPPgdp the data is closer to each other as expected with the transformation and we don't see Norway or Switzerland as potential as we did before.

Fig 8. Added-Variable Plots for 6 Predictors of 'modern_lm2' Model



9. Start by finding the best transformation of the response and then find transformations of the predictors. Do you end up with a different model than in 8?

Yes we would have a slightly different model if we start with a transformation for the response that is close to the MLE of $\lambda = 0.78$ from `boxCox` and is also interpretable (i.e. $\lambda = 0.5$). Assuming we proceed with this, we would find a λ for **Fertility** that is significantly different from 1 using the `boxTidwell` formula - $\lambda = 1.5$. To make it more interpretable I would have chosen $\lambda = 2$ instead. However, given the 95% confidence interval for the response λ includes 1 (as seen below), I wouldn't feel comfortable choosing any transformation in the first place. Thus, we would be back looking to find predictor λ s with an untransformed response.



10. Are there any outliers or influential points in the data? Explain. If so, refit the model after removing any outliers and comment on residual plots.

I used the Bonferroni Correction to test for outliers where $p_i < \frac{\alpha}{n}$ where $\alpha = 0.5$. From this test no observations have significant p-value [results and code hidden], thus none can be discarded as outliers. Regarding influential points, we can observe from **Fig 7** that even though with the transformations there are points with high leverage such as Cook Islands, Vanuatu and Kuwait. However, none are influential points.

```
pval = 2*(1 - pt(abs(rstudent(modern_lm2)), modern_lm2$df - 1))
rownames(na.omit(UN3_log))[pval < .05/nrow(na.omit(UN3_log))]
```

Summary of Results

11. For your final model, provide summaries of coefficients with 95% confidence intervals in a nice table with interpretations of each coefficient. These should be in terms of the original units!

Before finalizing the model and deciding which predictors to include, I tested whether keeping **Purban** would improve the model vs. a model without **Purban**. To do so I used an ANOVA to test H_0 - i.e. whether $\beta_6 = 0$ or not. From the resulting p-value, we fail to reject H_0 and therefore have excluded **Purban** from the final model,

$$\text{ModernC} = \beta_0 + \beta_1 \text{Change} + \beta_2 \log \text{PPgdp} + \beta_3 \text{Frate} + \beta_4 \log \text{Pop} + \beta_5 \text{Frate}$$

```
## Analysis of Variance Table
##
## Model 1: ModernC ~ Change + logPPgdp + Frate + logPop + Fertility
## Model 2: ModernC ~ Change + logPPgdp + Frate + logPop + Fertility + Purban
##   Res.Df  RSS Df Sum of Sq    F Pr(>F)
## 1     119 21420
## 2     118 21325   1    95.016 0.5258 0.4698
```

Below I've provided a summary table of the coefficients with 95% confidence interval in terms of the original units.

	2.5%	97.5%
(Intercept)	-24.569	32.773
Change	0.673	8.723
PPgdp	15.129	1098.940
Frate	0.050	0.349
Pop	1.223	14.598
Fertility	-12.595	-5.962

12. Provide a paragraph summarizing your final model and findings suitable for the US envoy to the UN after adjusting for outliers or influential points. You should provide a justification for any case deletions in your final model

Methodology

13. Prove that the intercept in the added variable scatter plot will always be zero. _Hint: use the fact that if H is the project matrix which contains a column of ones, then $1_n^T(I - H) = 0$. Use this to show that the sample mean of residuals will always be zero if there is an intercept.

Before starting some clarifications on the notation:

- $\hat{Y}_{(j)}$ refers to the Y vector regressed on all the predictors minus the j^{th} predictor
- $\hat{X}_{j(j)}$ refers the j^{th} predictor being regressed on the rest of the predictors (i.e. we regress the j^{th} predictor as if it were a response variable on the X matrix without this j^{th} predictor)
- Therefore, $X_j - \hat{X}_{j(j)} = \hat{e}_{j(j)}$ refers to the residuals from the regression of predictor j on the X matrix without the j predictor.
- Similarly, $Y - \hat{Y}_{(j)} = \hat{e}_{(j)}$ refers to the residuals from the regression the Y vector on the X matrix without the j predictor
- Although the correct notation for the hat matrix would be $H_{(j)}$, we will use H for simplicity given both regressions use the same hat matrix

Knowing the above we can prove that β_0 of an added variable plot will always be zero from,

$$\hat{e}_{(j)} = \vec{1}\hat{\beta}_0 + \hat{\beta}_j\hat{e}_{j(j)}$$

which takes the form similar to $Y = \hat{\beta}_0 + \hat{\beta}_1X_1$. Thus we can find $\hat{\beta}_j$ using the $(X^TX)^{-1}X^TY$ notation, but now using $X \equiv \hat{e}_{j(j)}$ and $Y \equiv \hat{e}_{(j)}$.

$$\begin{aligned}\hat{e}_{(j)} &= \vec{1}\hat{\beta}_0 + \overbrace{\left[((I - H)X_j)^T (I - H)X_j \right]^{-1} \left[(I - H)X_j \right]^T Y (I - H)X_j}^{\hat{\beta}_j} \\ &= \vec{1}\hat{\beta}_0 + \left[X_j^T \underbrace{(I - H)(I - H)}_{(I-H)} X_j \right]^{-1} \underbrace{X_j^T (I - H)Y}_{1 \times 1 \text{ scalar}} \underbrace{(I - H)X_j}_{1 \times 1 \text{ scalar}}\end{aligned}$$

If you multiply both sides by X_j^T and rearrange the scalars you get,

$$X_j^T(I-H)Y = X_j^T \vec{1} \hat{\beta}_0 + \underbrace{X_j^T(I-H)X_j [X_j^T(I-H)X_j]^{-1}}_I X_j^T(I-H)Y$$

$$\therefore X_j^T(I-H)Y = \sum_{i=1}^n X_{i,j} \hat{\beta}_0 + X_j^T(I-H)Y$$

$$\sum_{i=1}^n X_{i,j} \hat{\beta}_0 = X_j^T(I-H)Y - X_j^T(I-H)Y = 0$$

Thus we can see that the only way this relationship can only be zero is if $\hat{\beta}_0 = 0$ (i.e. the intercept is 0).

14. For multiple regression with more than 2 predictors, say a full model given by $Y \sim X_1 + X_2 + \dots + X_p$ we create the added variable plot for variable j by regressing Y on all of the X 's except X_j to form e_Y and then regressing X_j on all of the other X 's to form e_X . Confirm that the slope in a manually constructed added variable plot for one of the predictors in Ex. 10 is the same as the estimate from your model.

	Original	A-V Plot LM
Coeffs	-0.070768	-0.070768

